A Literate Program for Converting Tables to LongForm Dataframes

Jimmy Oh

Department of Statistics University of Auckland

Abstract

TableToLongForm automatically converts hierarchical Tables intended for a human reader into a simple LongForm Dataframe that is machine readable. It does this by recognising positional queues present in the hierarchical Table (which would normally be interpreted visually by the human brain) to decompose, then reconstruct the data into a LongForm Dataframe. This document provides a gallery of all recognised patterns and structures, with accompanying toy examples, before finally going into depth on the workings of the code itself.

Contents

1	Intr	roduction	3
	1.1	Motivation	3
	1.2	The Plan of Attack	4
	1.3	Recognised Patterns	4
		1.3.1 By Empty Below	5
		1.3.2 By Empty Below Transposed	5
			5
			6
			6
			7
			7
			7
		· · · · · · · · · · · · · · · · · · ·	8
2	Cod	de Overview	8
	2.1	Front End	9
	2.2		0
3	Idei	ntification 1	1
_	3.1	Identification - Main Function	1
			2
		v v	3
	3.2	1	4
			5
		1 0	5
			6

4	Disc	cern P	arentage	16
		4.0.4	plist	17
	4.1	Parent	tage - Front End Function	17
		4.1.1	Pare Col	17
	4.2		tage - Main Function	18
	4.3	Parent	tage - Low Level Functions	20
		4.3.1	Pare By Empty Right	20
		4.3.2	Pare By Empty Below	22
5	Rec	onstru	action	22
	5.1	Recon	struction - Main Function	23
	5.2	Recon	struction - Low Level Functions	25
		5.2.1	Reconstruction - Row Labels	
		5.2.2	Reconstruction - Column Labels	26
6	Chu	ınk Inc	dex	28
7	Idei	ntifier	Index	28
\mathbf{R}_{0}	efere	nces		30

On Literate Programs

This software is presented as a *literate program* written in the *noweb* format (Ramsey 1994). It serves as both the documentation and container of the literate program. The **noweb** file can be used to produce both the *literate document* and the executable code.

The literate document is separated into documentation chunks and named code chunks. Each code chunk can contain code directly, or contain references to other code chunks which act as placeholders for the contents of the respective code chunk. The name of each code chunk should serve as a short description of the code it contains. Thus each code chunk provides an overview of its purpose by either directly containing code, or by containing the names of other code chunks. The reader is then free to delve deeper into the respective code chunks if desired.

1 Introduction

1.1 Motivation

In recent times there has been a movement toward *Open Data*, particularly for government data^{1,2}, yet there is still a prevalence of data releases being for direct human consumption, rather than for machine consumption. One symptom of this is the release of data in tabular form that relies on the human ability to identify patterns and discern structure, in order to decipher the data (henceforth referred to as a Table). Such tables are difficult to read and analyse with the computer, signficantly limiting potential applications of this 'open' data.

LongForm is a simple alternative method of releasing the data that, due to its simplicity, is both easy to implement and is machine readable, greatly enhancing potential applications of the data. It is easy to go from a simple format such as a LongForm Dataframe to any number of other forms of presentation, including hierarchical Tables more suitable for direct human consumption. However the converse is rarely true. This is where TableToLongForm comes in, providing a way to automatically convert hierarchical Tables to a simple LongForm Dataframe, thus enabling much greater utilisation of the data.

	1	2	3	4	5	6	7	8	9	10	11	12	13
1	Labour Ford	e Status by	Sex by Sing	/Comb Ethni	c Group (Qr	tlv-Mar/Jun/	(Sep/Dec)						
2		Male											
3		European C	Dnlv							Maori Only			
4				Not in Labo	Working Ac	Labour Ford	Unemploym	Employmen	Total Labou	Persons Em	Persons Un	Not in Labo	Working Ag
5	2007Q4	855.8	20.0	280.0	1155.8	75.8	2.3	74.0	875.8	71.1	6.1	28.1	105.3
6	2008Q1	863.0	25.4	283.5	1171.9	75.8	2.9	73.6	888.5	69.1	7.5	31.4	107.9
7	2008Q2	850.1	26.0	280.7	1156.8	75.7	3.0	73.5	876.1	67.2	5.7	27.4	100.2
8	2008Q3	839.6	29.8	285.9	1155.3	75.2	3.4	72.7	869.4	71.7	8.7	30.7	111.1
9	2008Q4	854.8	29.5	274.7	1158.9	76.3	3.3	73.8	884.2	76.1	8.5	28.5	113.1
10	2009Q1	845.0	35.4	279.4	1159.8	75.9	4.0	72.9	880.4	75.4	8.4	35.7	119.5
11	2009Q2	831.6	34.9	279.7	1146.2	75.6	4.0	72.6	866.5	74.2	9.9	33.1	117.3
12	2009Q3	813.3	42.5	290.4	1146.2	74.7	5.0	71.0	855.8	70.9	10.9	36.0	117.8
13	2009Q4	831.1	40.1	277.0	1148.2	75.9	4.6	72.4	871.2	71.7	13.6	33.2	118.5
14	2010Q1	822.5	36.4	283.2	1142.1	75.2	4.2	72.0	858.9	71.8	11.3	35.3	118.4
15	2010Q2	825.3	39.9	290.1	1155.3	74.9	4.6	71.4	865.2	71.9	13.7	33.7	119.2
	2010Q3	836.9	31.0	287.1	1155.1	75.1	3.6	72.5	867.9	69.8	13.5	34.1	117.3
17	2010Q4	838.1	39.6	277.1	1154.8	76.0	4.5	72.6	877.7	70.7	14.4	36.4	121.5
18	2011Q1	829.6	36.8	281.2	1147.6	75.5	4.2	72.3	866.4	70.7	13.9	35.3	119.8
19	2011Q2	838.7		279.0	1158.7	75.9	4.7	72.4	879.6	67.1	10.5	37.1	114.7
20	2011Q3	830.5	34.6	280.2	1145.3	75.5	4.0	72.5	865.1	69.5	13.4	34.9	117.8
21	2011Q4	841.8	34.8	277.5	1154.1	76.0	4.0	72.9	876.6	69.2	12.7	36.1	118.0
22	2012Q1	843.1	43.3	282.7	1169.1	75.8	4.9	72.1	886.3	71.5	11.2	34.6	117.3
23	2012Q2	837.1	38.2	296.5	1171.8	74.7	4.4	71.4	875.3	66.2	10.6	33.0	109.7
24	2012Q3	833.0	38.0	298.4	1169.3	74.5	4.4	71.2	871.0	67.1	13.1	33.0	113.2
25	2012Q4	833.4	41.0	298.2	1172.6	74.6	4.7	71.1	874.4	63.0	12.3	35.4	110.7
26	2013Q1	832.0	35.8	294.7	1162.5	74.6	4.1	71.6	867.8	69.9	12.2	38.2	120.3
27	Table inforn	nation:											
28	Units:												
29	Persons En	nployed in La	abour Force:	Number, Ma	gnitude = T	housands							
30	Persons Un	employed in	Labour Ford	ce: Number,	Magnitude =	Thousands							
31	Not in Labo	ur Force: Nu	imber, Magn	itude = Thou	sands								
32	Working Ag	e Population	: Number, N	1agnitude = 7	Thousands								
33	Labour Ford	ce Participat	ion Rate: Pe	rcent, Magni	tude = Units								
34	Unemploym	ent Rate: Pe	ercent, Magr	nitude = Unit	3								
35	Employmen	t Rate: Perc	ent, Magnitu	ide = Units									
36	Total Labou	r Force: Nur	nber, Magnit	ude = Thous	ands								
37	Footnotes:												
38													
39	Symbols:												
40	figure not	available											
41	C: Confider	ntial											
42	E: Early Es	timate											
43	P: Provision	nal											

Figure 1: An example of a hierarchical Table. The data is Labour Force Status Survey data from Infoshare, Statistics New Zealand (2013) and in total spans 240 columns, making it suitable for neither man nor machine. However it is relatively tame in terms of how unsuitable for machines Tables can get and so could, with some manual labour, be read by a machine.

 $^{^1}$ "In many countries across the world, discussions, policies and developments are actively emerging around open access to government data." Davies and Bawa (2012a)

² "Over 100 OGD [Open Government Data] initiatives are active across the globe, ranging from community-led OGD projects in urban India, to a World Bank sponsored OGD programme in Kenya, government-led developments in Brazil, civil-society initiated work in Russia, and a World Wide Web Foundation supported programme in Ghana." Davies and Bawa (2012b)

	1	2	3	4	5	6	7	8	9	10	11	12	13
1				Persons En	Persons Un	Not in Labo	Working Ag	Labour Ford	Unemploym	Employmer	Total Labou	r Force	
2	Male	European (2007Q4	855.8	20	280	1155.8	75.8	2.3	74	875.8		
3	Male	European (2008Q1	863	25.4	283.5	1171.9	75.8	2.9	73.6	888.5		
4	Male	European (2008Q2	850.1	26	280.7	1156.8	75.7	3	73.5	876.1		
5	Male	European (2008Q3	839.6	29.8	285.9	1155.3	75.2	3.4	72.7	869.4		
6	Male	European (2008Q4	854.8	29.5	274.7	1158.9	76.3	3.3	73.8	884.2		
7	Male	European (2009Q1	845	35.4	279.4	1159.8	75.9	4	72.9	880.4		
8	Male	European (2009Q2	831.6	34.9	279.7	1146.2	75.6	4	72.6	866.5		
9	Male	European (2009Q3	813.3	42.5	290.4	1146.2	74.7	5	71	855.8		
10	Male	European (2009Q4	831.1	40.1	277	1148.2	75.9	4.6	72.4	871.2		
11	Male	European (2010Q1	822.5	36.4	283.2	1142.1	75.2	4.2	72	858.9		
12	Male	European (2010Q2	825.3	39.9	290.1	1155.3	74.9	4.6	71.4	865.2		
13	Male	European (2010Q3	836.9	31	287.1	1155.1	75.1	3.6	72.5	867.9		
14	Male	European (838.1	39.6	277.1	1154.8		4.5	72.6	877.7		
15	Male	European (829.6	36.8	281.2	1147.6	75.5	4.2	72.3	866.4		
16	Male	European (838.7	41	279	1158.7		4.7	72.4	879.6		
17	Male	European (2011Q3	830.5	34.6	280.2	1145.3	75.5	4	72.5	865.1		
18	Male	European (841.8	34.8	277.5	1154.1	76	4	72.9	876.6		
19	Male	European (2012Q1	843.1	43.3	282.7	1169.1	75.8	4.9	72.1	886.3		
20	Male	European (837.1	38.2	296.5	1171.8	74.7	4.4	71.4	875.3		
21	Male	European (833	38	298.4	1169.3		4.4	71.2	871		
22	Male	European (833.4	41	298.2	1172.6		4.7	71.1	874.4		
23	Male	European (832	35.8	294.7	1162.5	74.6	4.1	71.6	867.8		
24	Male	Maori Only	2007Q4	71.1	6.1	28.1	105.3	73.4	7.9	67.6	77.2		
25	Male	Maori Only		69.1	7.5	31.4	107.9	71	9.7	64.1	76.6		
26	Male	Maori Only		67.2	5.7	27.4	100.2		7.8	67	72.8		
27	Male	Maori Only		71.7	8.7	30.7	111.1	72.3	10.8	64.5	80.3		
28	Male	Maori Only		76.1	8.5	28.5	113.1	74.8	10	67.3	84.5		
29	Male	Maori Only		75.4	8.4	35.7	119.5	70.1	10.1	63.1	83.8		
30	Male	Maori Only		74.2	9.9	33.1	117.3	71.8	11.8	63.3	84.2		
31	Male	Maori Only		70.9	10.9	36	117.8	69.5	13.4	60.2	81.8		
32	Male	Maori Only		71.7	13.6	33.2	118.5	71.9	15.9	60.5	85.3		
33	Male	Maori Only		71.8	11.3	35.3	118.4	70.2	13.6	60.6	83.1		
34	Male	Maori Only		71.9	13.7	33.7	119.2	71.8	16	60.3	85.6		
35	Male	Maori Only		69.8	13.5	34.1	117.3	71	16.2	59.5	83.3		
36	Male	Maori Only		70.7	14.4	36.4	121.5	70	16.9	58.2	85.1		
37	Male	Maori Only		70.7	13.9	35.3	119.8	70.6	16.4	59	84.6		
38	Male	Maori Only		67.1	10.5	37.1	114.7	67.6	13.5	58.5	77.6		
39	Male	Maori Only		69.5	13.4	34.9	117.8	70.4	16.1	59	82.9		
40	Male	Maori Only		69.2	12.7	36.1	118	69.4	15.5	58.6	81.8		
41	Male	Maori Only		71.5	11.2	34.6	117.3	70.5	13.6	60.9	82.7		
42	Male	Maori Only		66.2	10.6	33	109.7	69.9	13.8	60.3	76.8		
43	Male	Maori Only	2012Q3	67.1	13.1	33	113.2	70.8	16.4	59.2	80.2		

Figure 2: An example of a LongForm Dataframe. This is the Labour Force Status Survey data after automatic conversion with TableToLongForm. Most, if not all statistical software (including spreadsheet software like Excel) can read, manipulate and run analyses on this without any problems.

1.2 The Plan of Attack

Unless the Table is horrible beyond mortal imagination, it should have some kind of pattern, such that a human will be able to discern the structure and hence understand the data it represents. This code attempts to algorithmically search for such patterns, discern the structure, then reconstruct the data into a LongForm Dataframe. Refer to Section 1.3 for a full gallery of currently recognised patterns.

The task can be seen to consist of three phases:

- Phase One is Identification (Section 3), which involves identifying the rows and columns where the labels and the data can be found.
- Phase Two is Discerning the Parentage (Section 4), which involves identifying the hierarchical structure of the data, based on the row and column labels.
- Phase Three is Reconstruction (Section 5), where we use what we've found in the first two phases to reconstruct the data into a LongForm Dataframe.

1.3 Recognised Patterns

Here we list, with toy examples, all the recognised patterns and structures. TableToLongForm should be able to process any combination of these patterns to automatically convert many different types of tables. The last example is a *complete* example that contains most of the recognised patterns in a single horrible table.

1.3.1 By Empty Below

	1	2	3	4	5	6
1			Column 1	Column 2	Column 3	Column 4
2	Row Parent1	Row Child1	10	20	30	40
3		Row Child2	11	21	31	41
4	Row Parent2	Row Child1	12	22	32	42
5		Row Child2	13	23	33	43

The most simple type of parentage, here the *parent* and *children* are in different columns and we can see which of the children belong to which parent through the use of empty space below each parent.

	1	2	3	4	5	6
1			Column 1	Column 2	Column 3	Column 4
2	Row Parent1	Row Child1	10	20	30	40
3	Row Parent1	Row Child2	11	21	31	41
4	Row Parent2	Row Child1	12	22	32	42
5	Row Parent2	Row Child2	13	23	33	43

1.3.2 By Empty Below Transposed

	1	2	3	4	5
1		Row Parent1		Row Parent2	
2		Row Child1	Row Child2	Row Child1	Row Child2
3	Column 1	10	11	12	13
4	Column 2	20	21	22	23
5	Column 3	30	31	32	33
6	Column 4	40	41	42	43

We note that parentage patterns recognised for row labels can often be applied to the transpose of column labels. This is how TableToLongForm deciphers most column labels, with some exceptions.

	1	2	3	4	5
1			Row Child1	Row Child2	
2	Row Parent1	Column 1	10	11	
3	Row Parent1	Column 2	20	21	
4	Row Parent1	Column 3	30	31	
5	Row Parent1	Column 4	40	41	
6	Row Parent2	Column 1	12	13	
7	Row Parent2	Column 2	22	23	
8	Row Parent2	Column 3	32	33	
9	Row Parent2	Column 4	42	43	

1.3.3 By Empty Right 1

	1	2	3	4	5	6	7
1				Column 1	Column 2	Column 3	Column 4
2	Row Parent1			10	20	30	40
3	Row Child1	Row Child-Ch	nild1	11	21	31	41
4	Row Child2	Row Child-Ch	nild2	12	22	32	42
5	Row Parent2			13	23	33	43
6	Row Child1	Row Child-Ch	nild1	14	24	34	44
7		Row Child-Ch	nild2	15	25	35	45

In this situation we have children in the same column as their parent. We can still recognise these as children if the children have children (*Child-Child*) in a different column, while the

parent does not (and hence is Empty Right).

Note the values pertaining to the Parent (if any) are discarded. This is because they are assumed to simply represent the sum of their children's values. It is planned for a sum-check to be implemented later to make this more robust.

	1	2	3	4	5	6	7
1				Column 1	Column 2	Column 3	Column 4
2	Row Parent1	Row Child1	Row Child-Ch	11	21	31	41
3	Row Parent1	Row Child2	Row Child-Ch	12	22	32	42
4	Row Parent2	Row Child1	Row Child-Ch	14	24	34	44
5	Row Parent2	Row Child1	Row Child-Ch	15	25	35	45

1.3.4 By Empty Right 2

	1	2	3	4	5	6
1			Column 1	Column 2	Column 3	Column 4
2	Row Parent1		10	20	30	40
3		Row Child1	11	21	31	41
4		Row Child2	12	22	32	42
5	Row Parent2		13	23	33	43
6		Row Child1	14	24	34	44
7		Row Child2	15	25	35	45

Here we have both Empty Below and Empty Right. Either algorithm can handle this situation, but simply due to the ordering of the algorithms such situations are handled as Empty Right.

	1	2	3	4	5	6
1			Column 1	Column 2	Column 3	Column 4
2	Row Parent1	Row Child1	11	21	31	41
3	Row Parent1	Row Child2	12	22	32	42
4	Row Parent2	Row Child1	14	24	34	44
5	Row Parent2	Row Child2	15	25	35	45

1.3.5 By Empty Right 3

	1	2	3	4	5	6	7	8
1				Column 1	Column 2	Column 3	Column 4	
2	Row Super-Page 1	arent1		10	20	30	40	
3	Row Parent1			11	21	31	41	
4	Row Child1	Row Child-Ch	ild1	12	22	32	42	
5	Row Parent2			13	23	33	43	
6	Row Child1	Row Child-Ch	ild1	14	24	34	44	
7	Row Super-Page 1	arent2		15	25	35	45	
8	Row Parent1			16	26	36	46	
9	Row Child1	Row Child-Ch	ild1	17	27	37	47	
10	Row Parent2			18	28	38	48	
11	Row Child1	Row Child-Ch	ild1	19	29	39	49	

The "parent-child in the same column" situation can be extended further. Here we have parents (Super-Parent) who have children (Parent), who each further have children (Child), all in the same column. Such situations can still be recognised if the lowest-level children in the column (Child) have children in a different column (Child-Child), while its direct parents (Parent) each have children in the same column (Child) but not in a different column (is Empty Right), and the top-most parents (Super-Parents) also have no children in a different column (is also Empty Right).

The algorithm cannot currently handle super-super-parents.

	1	2	3	4	5	6	7	8
1					Column 1	Column 2	Column 3	Column 4
2	Row Super-Page 1	Row Parent1	Row Child1	Row Child-Ch	12	22	32	42
3	Row Super-Page 1	Row Parent2	Row Child1	Row Child-Ch	14	24	34	44
4	Row Super-P	Row Parent1	Row Child1	Row Child-Ch	17	27	37	47
5	Row Super-Page 1	Row Parent2	Row Child1	Row Child-Ch	19	29	39	49

1.3.6 Multi-row Column Label

	1	2	3	4	5
1		Column	Column	Column	Column
2		Child1	Child2	Child3	Child4
3	Row 1	10	20	30	40
4	Row 2	11	21	31	41
5	Row 3	12	22	32	42
6	Row 4	13	23	33	43

Often column labels are physically split over multiple rows rather than making use of line breaks in the same cell. In such occurences, any row not identified as a parent are collapsed into a single row of labels. It is eventually planned for pattern recognition to be used here to make this collapsing smarter.

	1	2	3	4	5	
1		Column Child	Column Child2	Column Child3	Column Child	
2	Row 1	10	20	30	40	
3	Row 2	11	21	31	41	
4	Row 3	12	22	32	42	
5	Row 4	13	23	33	43	

1.3.7 Misaligned Column Label

	1	2	3	4	5	6	7	8	9
1			Column Parer	nt1			Column Parer		
2		Col Child1	Col Child2	Col Child3	Col Child4	Col Child1	Col Child2	Col Child3	Col Child4
3	Row 1	10	20	30	40	50	60	70	80
4	Row 2	11	21	31	41	51	61	71	81
5	Row 3	12	22	32	42	52	62	72	82
6	Row 4	13	23	33	43	53	63	73	83

Often column parents are physically centred over their children (N.B. where a spreadsheet's cell-merge feature is used to do the centering, the actual value is usually stored in the top-left cell and hence causes no problems). TableToLongForm makes use of pattern recognition to identify repeating patterns in the labels of the children, to help discern the correct parent for the children.

	1	2	3	4	5	6	7	8	9
1			Col Child1	Col Child2	Col Child3	Col Child4			
2	Column Paren	Row 1	10	20	30	40			
3	Column Paren	Row 2	11	21	31	41			
4	Column Paren	Row 3	12	22	32	42			
5	Column Paren	Row 4	13	23	33	43			
6	Column Paren	Row 1	50	60	70	80			
7	Column Paren	Row 2	51	61	71	81			
8	Column Paren	Row 3	52	62	72	82			
9	Column Paren	Row 4	53	63	73	83			

1.3.8 Find Single Table

	1	2	3	4	5
1	MISC INFORM	MATION			
2	MISC INFORM	MATION			
3		Column 1	Column 2	Column 3	Column 4
4	Row 1	10	20	30	40
5	Row 2	11	21	31	41
6	Row 3	12	22	32	42
7	Row 4	13	23	33	43
8	MISC INFORM	MATION	MISC INFORM		
9	MISC INFORM	MATION	MISC INFORM		

A table is often found amongst miscellaneous information we do not want. TableToLongForm is intended to have several algorithms to identify not only a single table, but multiple tables on the same 'page'. Currently however, it can only identify a single table per 'page' by searching for a block (rectangular region) of numbers, which is assumed to be our table of data

	1	2	3	4	5
1		Column 1	Column 2	Column 3	Column 4
2	Row 1	10	20	30	40
3	Row 2	11	21	31	41
4	Row 3	12	22	32	42
5	Row 4	13	23	33	43

1.3.9 Complete Example

	1	2	3	4	5	6	7	8	9	10	11
1	MISC INFORM	MATION									
2					Column Par	ent1			Column Par	ent2	
3				Column	Column	Column	Column	Column	Column	Column	Column
4				Child1	Child2	Child3	Child4	Child1	Child2	Child3	Child4
5	Row Super-P	arent		10	20	30	40	50	60	70	80
6	Row Parent1			11	21	31	41	51	61	71	81
7	Row Child1	Row Child-Ch	nild1	12	22	32	42	52	62	72	82
8		Row Child-Ch	nild2	13	23	33	43	53	63	73	83
9	Row Child2	Row Child-Ch	nild1	14	24	34	44	54	64	74	84
10		Row Child-Ch	nild2	15	25	35	45	55	65	75	85
11	Row Parent2			16	26	36	46	56	66	76	86
12	Row Child1	Row Child-Ch	nild1	17	27	37	47	57	67	77	87
13		Row Child-Ch	nild2	18	28	38	48	58	68	78	88
14	Row Child2	Row Child-Ch	nild2	19	29	39	49	59	69	79	89
15	MISC INFORM	MATION									
16	MISC INFORM	MATION									

A complete example containing a combination of many of the patterns listed above.

	1	2	3	4	5	6	7	8	9	10	11
1						Column Child	Column Child:	Column Child	Column Child4	1	
2	Column Paren	Row Super-P	Row Parent1	Row Child1	Row Child-Ch	12	22	32	42		
3	Column Paren	Row Super-P	Row Parent1	Row Child1	Row Child-Ch	13	23	33	43		
4	Column Paren	Row Super-P	Row Parent1	Row Child2	Row Child-Ch	14	24	34	44		
5	Column Paren	Row Super-P	Row Parent1	Row Child2	Row Child-Ch	15	25	35	45		
6	Column Paren	Row Super-P	Row Parent2	Row Child1	Row Child-Ch	17	27	37	47		
7	Column Paren	Row Super-P	Row Parent2	Row Child1	Row Child-Ch	18	28	38	48		
8	Column Paren	Row Super-P	Row Parent2	Row Child2	Row Child-Ch	19	29	39	49		
9	Column Paren	Row Super-P	Row Parent1	Row Child1	Row Child-Ch	52	62	72	82		
10	Column Paren	Row Super-P	Row Parent1	Row Child1	Row Child-Ch	53	63	73	83		
11	Column Paren	Row Super-P	Row Parent1	Row Child2	Row Child-Ch	54	64	74	84		
12	Column Paren	Row Super-P	Row Parent1	Row Child2	Row Child-Ch	55	65	75	85		
13	Column Paren	Row Super-P	Row Parent2	Row Child1	Row Child-Ch	57	67	77	87		
14	Column Paren	Row Super-P	Row Parent2	Row Child1	Row Child-Ch	58	68	78	88		
15	Column Paren	Row Super-P	Row Parent2	Row Child2	Row Child-Ch	59	69	79	89		

2 Code Overview

TableToLongForm is structured as follows.

```
8 \langle Table To Long Form. R \ 8 \rangle \equiv
\langle document \ header \ 9a \rangle
\langle Front \ End \ 9b \rangle
\langle Identification \ 11a \rangle
\langle Discern \ Parentage \ 16b \rangle
\langle Reconstruction \ 22b \rangle
\langle Back \ End \ 10 \rangle
```

This code is written to file ${\tt TableToLongForm.R.}$

We place a document header at the top of the extracted code to encourage people to read the literate description rather than attempting to study the code alone.

 $\langle document \ header \ 9a \rangle \equiv$

9a

2.1 Front End

The main function TableToLongForm is defined here. For most users this is the only function they will call. However, the majority of the supporting functions are not hidden and therefore can easily be viewed and/or modified by users.

```
⟨Front End 9b⟩≡
9b
          TableToLongForm =
            function(datamat, IdentResult = NULL,
                     fulloutput = FALSE, diag = FALSE, diagname = NULL){
                if(is.null(diagname)) diagname = deparse(substitute(datamat))
                assign("TCRunout", file(paste0(diagname, ".TCRunout"),
                                         "w"), envir = .GlobalEnv)
                on.exit({
                  close(TCRunout)
                  rm("TCRunout", envir = .GlobalEnv)
                })
              }
              fullout = ReconsMain(datamat, IdentResult)
              if(fulloutput) fullout else fullout$datafr
            }
```

2.2 Back End

10

Various code, mainly to help produce diagnostic output, can be ignored by most users.

print.plist A print method for class plist, which are nested lists with a numeric vector at the lowest level; **print.default** is rather inefficient in displaying such nested lists.

TCRsink Sinks the output to TCRunout for diagnostic output. Requires the existence of TCRunout which is created by the main function TableToLongForm when diag = TRUE.

Spaces may be introduced by match.call, thus any spaces in the args of variables to sink (that is, the arguments supplied via ...) are removed without warning.

```
\langle Back\ End\ 10 \rangle \equiv
    print.plist = function(plist){
      plistC = function(plist){
        pLoc = attr(plist, "Loc")
         if(is.list(plist)){
          namevec = names(plist)
           if(!is.null(pLoc))
             namevec = pasteO(names(plist),
               " (", pLoc[,"rows"], ", ", pLoc[,"cols"], ")")
          namelist = as.list(namevec)
           for(i in 1:length(namelist))
             namelist[[i]] =
               c(paste("+", namelist[[i]]),
                 paste("-", plistC(plist[[i]])))
           do.call(c, namelist)
         } else{
           if(!is.null(names(plist))){
             namevec = names(plist)
             if(!is.null(pLoc))
               namevec = pasteO(names(plist),
                 " (", plist, ", ", pLoc[,"cols"], ")")
             paste("+", namevec)
           } else paste(plist, collapse = " ")
         }
       cat(plistC(plist), sep = "\n")
    attrLoc =
      function(plist, rows = NULL, cols = NULL){
         attr(plist, "Loc") = cbind(rows, cols)
         class(plist) = "plist"
        plist
      }
    TCRsink =
      function(ID, ...)
      if(exists("TCRunout", envir = .GlobalEnv)){
        varlist = list(...)
        names(varlist) = gsub(" ", "", as.character(match.call()[-(1:2)]))
        sink(TCRunout)
        for(i in 1:length(varlist)){
           cat("###TCR", ID, names(varlist)[i], "\n")
          print(varlist[[i]])
         sink()
```

```
}
Defines:
attrLoc, used in chunks 18-22.
TCRsink, used in chunks 11-14 and 18-24.
```

3 Identification

We separate the Identification functions into two groups.

Ident Main contains the main function that is called by the *Front End* function.

Ident Low Level contains supporting functions called by the *Ident Main* function.

```
11a \langle Identification \ 11a \rangle \equiv \langle Ident \ Main \ 11b \rangle \langle Ident \ Low \ Level \ 14d \rangle
```

3.1 Identification - Main Function

The purpose of the IdentMain function is to identify where in the file the data is found and where the accompanying labels are, while ignoring any extraneous information we do not want. It should also identify the presence of multiple tables in the same file.

It is intended for this procedure to involve a number of Identification algorithms that are used for a high degree of reliability and flexibility, but at this stage there is only a single algorithm.

The algorithms used are:

• Ident by Most Common Boundary.

Algorithms planned for the near future are:

• Ident by Runs.

The output of IdentMain will be a list containing two elements, rows and cols, each of which is a list containing these two elements:

label - a vector of the rows or columns where the labels are found.

data - a vector of the rows or columns where the data are found.

```
Ident Main 11b⟩≡

IdentMain =

function(datamat){

⟨Ident by Most Common Boundary 12a⟩

⟨Group Column Labels 13b⟩

TCRsink("IM", rowslist, colslist)

list(rows = rowslist, cols = colslist)

}

Defines:

IdentMain, used in chunk 23b.

Uses TCRsink 10.
```

```
> rowslist
$label
[1] 1 2 3 4

$data
   [1] 5 6 7 8 9 10 11 12 13 14

> colslist
$label
[1] 1 2

$data
$data[1]
[1] 4 5 6 7

$data[2]
[1] 8 9 10 11
```

3.1.1 Ident By Most Common Boundary

The IdentMostCommonBoundary Low Level function is used to find the most common start and end rows and columns (the boundary) to search for a block (rectangular region) of numbers, which is assumed to be our table of data.

```
⟨Ident by Most Common Boundary 12a⟩≡
12a
             ⟨Get Non empty rows and cols 12b⟩
             \langle Call\ Ident\ MostCommonBoundary\ 12c \rangle
             ⟨Construct rowslist and colslist 13a⟩
12b
        \langle Get \ Non \ empty \ rows \ and \ cols \ 12b \rangle \equiv
             rowNonempty = (1:nrow(datamat))[IdentNonEmpty(datamat, 1)]
             colNonempty = (1:ncol(datamat))[IdentNonEmpty(datamat, 2)]
        Uses \ {\tt IdentNonEmpty} \ 15a.
        \langle Call\ Ident\ MostCommonBoundary\ 12c \rangle \equiv
12c
             rowData = IdentMostCommonBoundary(datamat, 2)
             colData = IdentMostCommonBoundary(datamat, 1)
             ## Temporary fix for first col being all numbers (e.g. years)
             if(colData[1] == 1) colData[1] = 2
             TCRsink("CIMCB", rowData, colData)
        Uses IdentMostCommonBoundary 16a and TCRsink 10.
```

```
> rowData
[1] 5 14
> colData
[1] 4 11
```

We construct the interim rowslist taking every non-empty row before the most common start of the numbers block (rowData[1]) and assigning these to the label region. The numbers block (which is bounded by rowData[1] and rowData[2]) is assigned to the data region. The interim colslist is constructed in the same manner.

```
13a
       \langle Construct \ rowslist \ and \ colslist \ 13a \rangle \equiv
            rowslist = list(label = rowNonempty[rowNonempty < rowData[1]],</pre>
                              data = rowNonempty[(rowNonempty >= rowData[1]) &
                                                    (rowNonempty <= rowData[2])])</pre>
            colslist = list(label = colNonempty[colNonempty < colData[1]],</pre>
                              data = colNonempty[(colNonempty >= colData[1]) &
                                                    (colNonempty <= colData[2])])</pre>
            TCRsink("CRAC", rowslist, colslist)
       Uses TCRsink 10.
          Example values for ToyExComplete.csv
       > rowslist
       $label
       [1] 1 2 3 4
       $data
        [1] 5 6 7 8 9 10 11 12 13 14
       > colslist
       $label
       [1] 1 2
       $data
       [1] 4 5 6 7 8 9 10 11
```

3.1.2 Group Column Labels

We look for a repeating pattern in the column labels to handle cases of Misaligned Column Label (see Section 1.3.7).

```
13b \langle Group\ Column\ Labels\ 13b \rangle \equiv \langle Generate\ Pattern\ vector\ 14a \rangle \langle Take\ Largest\ Pattern\ 14b \rangle \langle Group\ by\ Pattern\ 14c \rangle
```

We loop through each row of the labels region and check for a pattern in either the contents of the cells (if they are all non-empty), or a pattern in which cells are empty (if any cells are empty), and store the results in Patvec (Pattern Vector).

Where multiple patterns are found, we assume the shorter patterns are patterns of families that are children to the parents of the largest pattern. Thus we always take the largest pattern

Some problems with this and next chunk, refer to BUG-ID 2.

```
14b \langle Take\ Largest\ Pattern\ 14b \rangle \equiv Patvec = max(Patvec)
```

If a pattern is found (NA = No patterns found), we group the columns into separate elements in a list and update the colslist. For easy handling later, irrespective of whether a pattern is found, colslist\$data is a list.

3.2 Identification - Low Level Functions

Here we discuss the low level functions that are called by the main Identification function. Each chunk corresponds to a separate low level function.

```
14d \langle Ident\ Low\ Level\ 14d \rangle \equiv
\langle Ident\ Non\ Empty\ 15a \rangle
\langle Ident\ Pattern\ 15b \rangle
\langle Ident\ Most\ Common\ Boundary\ 16a \rangle
```

3.2.1 IdentNonEmpty

Given a matrix (datamat) and a margin (1 for rows, 2 for columns), return a vector giving the indices of non-empty rows or columns. Can specify a different empty identifying function (default is.na). Procedure:

- 1. Compute isnonempty, a logical vector about whether the rows or cols are not empty.
- 2. Use which on isnonempty to get indices.

IdentNonEmpty, used in chunks 12b and 23c.

```
⟨Ident Non Empty 15a⟩≡
    IdentNonEmpty =
        function(datamat, margin, emptyident = is.na){
        isnonempty = apply(datamat, margin, function(x) !all(emptyident(x)))
        which(isnonempty)
    }
Defines:
```

3.2.2 IdentPattern

15a

Attempt to discern a repeating pattern in vec, which can be a vector of any type (which is coerced to character). The returned value, res is either NA if no pattern is found. Or it is the grouping number for the repeating pattern, e.g.

```
vec = 1 1 1 1, then res = 1
vec = 3 4 3 4, then res = 2
vec = 1 2 3 1, then res = NA
```

IdentPattern does this fairly efficiently by use of regular expressions. It combines the first i elements of vec and collapses this into a single string. A grep is then called on the entire vec that has been collapsed into a single string, checking to see if the entire string can be matched to some repeat of the aforementioned collapsed string of the first i elements.

For the moment it is possible for this to fail (and can even be intentionally gamed by providing something like vec = c(12, 1, 2), which will return a pattern of 1, when it should return NA), so it should be changed to be more reliable (though less efficient).

```
\langle Ident\ Pattern\ 15b \rangle \equiv
15b
            IdentPattern =
              function(vec){
                len = length(vec)
                res = NA
                for(i in 1:floor(len/2)){
                  curseg = paste("^(", paste(vec[1:i], collapse = ""),
                    ")+$", sep = "")
                  if(nchar(curseg) > 2559){
                    warning("Label lengths too long for regular expressions to ",
                             "work. IdentPattern has been aborted. A pattern may ",
                             "exist but it cannot be found with the current ",
                             "algorithm.")
                  } else if(length(grep(curseg, paste(vec, collapse = ""))) > 0){
                    res = i
                    break
                }
                res
              }
       Defines:
```

IdentPattern, used in chunk 14a.

3.2.3 Ident Most Common Boundary

Search for the most common first and last rows/cols to identify a block (rectangular region) of numbers. Procedure:

- 1. Suppose margin = 2, then loop through each column and search for cells containing numbers.
- 2. Compute the first row with a number for each column (nstarts), and do the same for the last row (nends).
- 3. Return the most common first and last rows.

IdentMostCommonBoundary, used in chunk 12c.

4 Discern Parentage

We separate the Parentage functions into three groups.

Pare Front is a simple 'front-end' function that makes the appropriate first call to PareMain, and is the function called by the *Front End* function.

Pare Col A specialised front-end to PareFront that handles various fringe cases for Discering Parentage for Column Labels, before eventually calling PareFront.

Pare Main contains the main function that recursively call itself until the all parentage is discerned.

Pare Low Level contains supporting functions called by the Pare Main function.

See the section on the main function (Section 4.2) for details on the purpose of the *Discern Parentage* stage.

```
16b \langle Discern\ Parentage\ 16b \rangle \equiv \langle Pare\ Front\ 17a \rangle \langle Pare\ Col\ 17b \rangle \langle Pare\ Main\ 18c \rangle \langle Pare\ Low\ Level\ 20b \rangle
```

4.0.4 plist

explanation.

4.1 Parentage - Front End Function

This front end function takes the datamat and constructs an initialising plist (Parentage List), which is used to make the first call to the main function.

```
17a \langle Pare\ Front\ 17a \rangle \equiv
PareFront =
function(datamat)
PareMain(datamat = datamat, plist =
list(rows = 1:nrow(datamat), cols = 1:ncol(datamat)))
Defines:
PareFront, used in chunks 18b and 23c.
Uses PareMain 18c.
```

4.1.1 Pare Col

The Parentage functions were initially designed to work with Row Labels only, however we can also use them to discern the parentage of Col Labels once we handle a few differences. We define a front-end to the front-end function called PareCol to do this.

```
17b \langle Pare\ Col\ 17b \rangle \equiv
PareCol =
function(datamat, datacols, labelrows){
\langle Case\ Misaligned\ Col\ Parent\ 17c \rangle
datacols = unlist(datacols)
\langle Collapse\ Fullrow\ Labels\ 18a \rangle
\langle Call\ Pare\ Front\ 18b \rangle
}
Defines:
PareCol, used in chunk 24.
Uses datacols 23b and labelrows 23b.
```

Unlike with Row Labels where the parents are reliably in the top-left corner of their family, Col Label parents are sometimes 'misaligned'. In some cases this arises as Col Label parents might be centred over their family. Other times, it happens for no apparent logical reason.

Regardless of the cause, we need to correct for this. During the Identification phase, we Grouped the Column Labels based on repeating patterns. We use these groupings to identify the families, and if the parent is not found where it should be, we simply shift it over to the right place.

Uses datacols 23b and labelrows 23b.

It is also quite common for Col Labels that are too wide to be physically split over multiple rows to manage the width of the labels. For now, we simply assume that any rows that are not full (and hence not parents) should all really be a single row of children, and collapse these.

```
⟨Collapse Fullrow Labels 18a⟩≡
18a
            notfullrows = apply(datamat[labelrows, datacols, drop = FALSE], 1,
              function(x) any(is.na(x)))
            if(any(diff(notfullrows) > 1))
              warning("full rows followed by not full rows!")
            pastestring = ""
            pasterows = which(!notfullrows)
            for(i in 1:length(pasterows))
              pastestring[i] = paste("datamat[labelrows[", pasterows[i],
                            "], datacols]", sep = "")
            collapsedlabels = eval(parse(text = paste("paste(",
                                             paste(pastestring, collapse = ", "),
                                             ")", sep = "")))
       Uses datacols 23b and labelrows 23b.
          Once the above is handled, we can simply transpose our Col Labels and call PareFront
       on it.
       ⟨Call Pare Front 18b⟩≡
18b
            labeldatamat = rbind(datamat[labelrows[notfullrows], datacols],
              collapsedlabels)
            PareFront(t(labeldatamat))
       Uses datacols 23b, labelrows 23b, and PareFront 17a.
```

4.2 Parentage - Main Function

The purpose of the PareMain function is to identify (or *Discern*, to better differentiate this stage from the *Identification* stage) hierarchical relationships (the *Parentage*) in the data.

It first makes various checks for fringe cases, then calls various detection algorithms (Pare Low Levels) to discern the parentage.

```
Pare Main 18c⟩≡

PareMain =

function(datamat, plist){

⟨If only one column 18d⟩

⟨If first column empty 19a⟩

⟨If only one row 19b⟩

⟨If first cell empty 19c⟩

⟨Otherwise call Pare Low Levels 20a⟩

class(res) = "plist"

res

}

Defines:

PareMain, used in chunks 17a, 19a, and 20a.
```

If only one column is found then this means we are in the right-most column (or there was only one column to begin with), and hence the currently examined cells cannot be parents. We return the rows of these children as a vector, with names that correspond to their labels.

Uses attrLoc 10 and TCRsink 10.

```
> plist
       $rows
       [1] 3 4
       $cols
       [1] 2
       > res
       + Row Child-Child1 (3, 2)
       + Row Child-Child2 (4, 2)
          If the first column is found to be empty, then we will shift to the next column (which we
       know exists because we passed the check for only one column).
19a
       \langle If\ first\ column\ empty\ 19a \rangle \equiv
            else if(all(is.na(datamat[plist$rows, plist$cols[1]]))){
               plist$cols = plist$cols[-1]
               res = PareMain(datamat, plist)
       Uses PareMain 18c.
          If only one row is found then our row is a parent to itself (we know there are children in
       the row as we passed the check for only one column). We return the row as a numeric vector,
       nested in a list using correct parentage and names of the parentage within the row.
19b
       \langle If \ only \ one \ row \ 19b \rangle \equiv
            else if(length(plist$rows) == 1){
               res = structure(plist$rows,
                 .Names = datamat[plist$rows, plist$cols[length(plist$cols)]])
               res = attrLoc(res, cols = plist$cols[length(plist$cols)])
               for(i in (length(plist$cols) - 1):1){
                 res = list(res)
                 names(res) = datamat[plist$rows, plist$cols[i]]
                 res = attrLoc(res, rows = plist$rows, cols = plist$cols[i])
               TCRsink("IOOR", plist, res)
            }
       Uses attrLoc 10 and TCRsink 10.
          Example values for ToyExComplete.csv
       Never occurs
          If the first cell is empty, after all previous checks, then this is an unrecognised format and
       we return a warning message.
       \langle If first \ cell \ empty \ 19c \rangle \equiv
19c
            else if(is.na(datamat[plist$rows[1], plist$cols[1]])){
               warning("cell[1, 1] is empty")
               print(plist)
               res = NA
            }
```

If we have passed all the checks, we can then call the Low Level Pare functions. We first call ByEmptyRight to check for *empty right* situations. If none are found, it returns NA, in which case we try ByEmptyBelow instead.

We then loop through each element of the returned list and call the main function, as per the recursive nature of the function.

```
20a  ⟨Otherwise call Pare Low Levels 20a⟩≡
    else{
        res = PareByEmptyRight(datamat, plist)
        if(any(is.na(res)))
        res = PareByEmptyBelow(datamat, plist)
        for(i in 1:length(res))
        res[[i]] = PareMain(datamat, res[[i]])
        res
}
```

Uses PareByEmptyBelow 22a, PareByEmptyRight 20c, and PareMain 18c.

4.3 Parentage - Low Level Functions

The Low Level Parentage functions are called by the Main Parentage function. In particular, ByEmptyRight is always called first. Then ByEmptyBelow is called on the results of the above.

```
20b \langle Pare\ Low\ Level\ 20b \rangle \equiv \langle Pare\ By\ Empty\ Right\ 20c \rangle \langle Pare\ By\ Empty\ Below\ 22a \rangle
```

20c

20d

4.3.1 Pare By Empty Right

We check to see if we have an empty right situation. If we do not, we return NA.

```
⟨Pare By Empty Right 20c⟩≡
PareByEmptyRight =
function(datamat, plist)
with(plist,
    if(all(is.na(datamat[rows[1], cols[-1]]))){
    ⟨Check for Other Empty Rights 20d⟩
    ⟨Case Single Empty Right 20e⟩
    ⟨Case Multiple Empty Rights 21⟩
    res
} else NA)
Defines:
PareByEmptyRight, used in chunk 20a.
⟨Check for Other Empty Rights 20d⟩≡
emptyrights = apply(datamat[rows, cols[-1], drop = FALSE], 1,
function(x) all(is.na(x)))
```

In the case of only a single empty right, we know there is only a single parent, which is the first line. Thus we take everything except the first line (which will be the rows of the children of this parent) and pass this through with correct naming.

rowemptyright = rows[emptyrights]

New Zealand	
Auckland	
Accounting	Male
	Female
Economics	Male
	Female
Statistics	Male
	Female
Wellington	
Economics	Male
	Female
Statistics	Male
	Female
Australia	
Sydney	
Accounting	Male
	Female
Economics	Male
	Female
	Auckland Accounting Economics Statistics Wellington Economics Statistics Australia Sydney Accounting

Consider the toy example on the left.

In this case we do not have a simple ByEmptyRight structure. We have *super-parents* in the form of countries (New Zealand and Australia), and also *parents* in the form of cities (Auckland, Wellington and Sydney). To handle situations such as this, we must Check for Other Empty Rights.

If only a **Single Empty Right** is found, the situation is simple and we simply pass on the children of the single parent for the next iteration of PareMain.

However, if Multiple Empty Rights are found, we must identify the super-parents, and pass on the *children* of these super-parents (which would, in turn, contain parents and their children) as a list, to be handled in the next iteration of PareMain. In this example, we would have a list of length 2. The first element of the list would contain the plist with rows 2 to 13 (corresponding to the children of the New Zealand super-parent). The second element would have rows 15 to 19.

Example values for ToyExComplete.csv

```
> res
Never occurs
```

Uses attrLoc 10 and TCRsink 10.

21

In the case of multiple empty rights, we first call diff to compute the gap in rows between the empty rights. If the value of rowdiff[i] is 1, this means there is no gap between the i^{th} rowemptyright and the (i+1) rowemptyright. This happens with *super-parents* as described in the example above. In this case, we gather these super-parents and ignore all other rowemptyright (the parents inside the super-parents will be handled at the next iteration of PareMain). Note, we assume there are never any super-super-parents (i.e. we can only handle a maximum of 2-levels of parentage in the same column).

Whether or not super-parents were identified, we compute the rows for the children of each parent (or super-parent) identified by rowemptyright and pass this through as a list, with correct naming.

```
else{
    rowdiff = diff(rowemptyright)
    if(any(rowdiff == 1))
        rowstart = pmin(rowemptyright = 1, max(rows))
    rowend = c(pmax(rowemptyright[-1] - 1, min(rows)), max(rows))

res = list()
    for(i in 1:length(rowstart))
        res[i] = list(list(rows = rowstart[i]:rowend[i], cols = cols))
        names(res) = datamat[rowemptyright, cols[1]]
    res = attrLoc(res, rows = rowemptyright, cols = cols[1])
        TCRsink("CMER", res)
}
```

21

```
> res
+ Row Super-Parent (1, 1)
- + rows
- - 2 3 4 5 6 7 8 9 10
- + cols
- - 1 2
```

4.3.2 Pare By Empty Below

We check which cells are empty below (there should be at least 1 based on previous checks). Based on this, we compute the rows for the children of each parent and pass this through as a list, with correct naming.

```
\langle Pare\ By\ Empty\ Below\ 22a \rangle \equiv
22a
            PareByEmptyBelow =
              function(datamat, plist)
              with(plist, {
                emptybelow = is.na(datamat[rows, cols[1]])
                rowstart = rows[!emptybelow]
                rowend = c(rowstart[-1] - 1, max(rows))
                res = list()
                for(i in 1:length(rowstart))
                  res[i] = list(list(rows = rowstart[i]:rowend[i], cols = cols[-1]))
                names(res) = datamat[rowstart, cols[1]]
                res = attrLoc(res, rows = rowstart, cols = cols[1])
                TCRsink("PBEB", res)
                res
              })
       Defines:
            PareByEmptyBelow, used in chunk 20a.
       Uses attrLoc 10 and TCRsink 10.
          Example values for ToyExComplete.csv
       > res
       + Row Child1 (3, 1)
       - + rows
       - - 3 4
       - + cols
       - - 2
       + Row Child2 (5, 1)
       - + rows
       - - 5 6
       - + cols
       - - 2
```

5 Reconstruction

We separate the Reconstruction functions into two groups.

Recons Main contains the main function that is called by the Front End function.

Recons Low Level contains supporting functions called by the Recons Main function.

```
22b \langle Reconstruction \ 22b \rangle \equiv \langle Recons \ Main \ 23a \rangle \langle Recons \ Low \ Level \ 25a \rangle
```

5.1 Reconstruction - Main Function

The ReconsMain function is, in a manner of speaking, the true TableToLongForm function, as it makes the calls to IdentMain and PareFront, in conjunction with its own Recons Low Level functions, to carry out the conversion.

```
23a \langle Recons\ Main\ 23a \rangle \equiv

ReconsMain =

function(datamat, IdentResult){

\langle Call\ Ident\ Main\ 23b \rangle

\langle Reconstruct\ Row\ Labels\ 23c \rangle

\langle Reconstruct\ Col\ Labels\ 24 \rangle
}
```

Call IdentMain and assign them meaningful names for convenience, the labelcols being the columns where the labels can be found, etc. These should all be a vector, except datacols which is a list.

```
if (is.null(IdentResult))

IdentResult = IdentMain(datamat)

labelcols = IdentResult$cols$label

datacols = IdentResult$cols$data

labelrows = IdentResult$rows$label

datarows = IdentResult$rows$label

datarows = IdentResult$rows$data

Defines:

labelcols, used in chunk 23c.

datacols, used in chunks 17, 18, and 24.

labelrows, used in chunks 17, 18, and 24.

datarows, used in chunks 23c and 24.

Uses IdentMain 11b.
```

We create a subset of datamat that contains just the Row Labels. We also remove any columns that are completely empty (N.B. this may no longer be necessary, need to do some testing).

We call PareFront on our subset to discern the parentage of the Row Labels. We then use this to reconstruct the portion of the LongForm Dataframe relating to the Row Labels and assign this to rowvecs.

```
23c \( \langle Reconstruct Row Labels 23c \rangle = \ \ \text{datamatRowLabels} = \text{datamat[datarows, labelcols, drop = FALSE]} \\ \text{datamatRowLabels} = \text{datamatRowLabels[,} \\ \text{IdentNonEmpty(datamatRowLabels, 2), drop = FALSE]} \\ \text{rowplist} = \text{PareFront(datamatRowLabels)} \\ \text{rowvecs} = \text{ReconsRowLabels(rowplist)} \\ \text{TCRsink("RRL", rowplist, rowvecs[1:4,])} \\ \text{Defines:} \\ \text{rowplist, used in chunk 24.} \\ \text{rowvecs, used in chunk 24-27.} \\ \text{Uses datarows 23b, IdentNonEmpty 15a, labelcols 23b, PareFront 17a, ReconsRowLabels 25b, and TCRsink 10.} \end{array}
```

```
> rowplist
+ Row Super-Parent (1, 1)
- + Row Parent1 (2, 1)
- - + Row Child1 (3, 1)
- - - + Row Child-Child1 (3, 2)
- - - + Row Child-Child2 (4, 2)
- - + Row Child2 (5, 1)
- - - + Row Child-Child1 (5, 2)
- - - + Row Child-Child2 (6, 2)
- + Row Parent2 (7, 1)
- - + Row Child1 (8, 1)
- - - + Row Child-Child1 (8, 2)
- - - + Row Child-Child2 (9, 2)
- - + Row Child2 (10, 1)
- - - + Row Child-Child2 (10, 2)
> rowvecs[1:4,]
 [,1]
                    [,2]
                                   [,3]
 "Row Super-Parent" "Row Parent1" "Row Child1" "Row Child-Child1"
 "Row Super-Parent" "Row Parent1" "Row Child1" "Row Child-Child2"
 "Row Super-Parent" "Row Parent1" "Row Child2" "Row Child-Child1"
 "Row Super-Parent" "Row Parent1" "Row Child2" "Row Child-Child2"
```

Due to the various fringe cases that exist with Col Labels, we cannot simply call PareFront and must instead call PareCol. We then create a subset of datamat that contains just the Col Labels and call ReconsColLabels which in truth reconstruct the entire LongForm Dataframe by making use of the rowvecs generated above.

Uses datacols 23b, datarows 23b, labelrows 23b, PareCol 17b, ReconsColLabels 26a, rowplist 23c, rowvecs 23c, and TCRsink 10.

```
> colplist
+ Column Parent1 (1, 2)
- + Column Child1 (1, 3)
- + Column Child2 (2, 3)
- + Column Child3 (3, 3)
- + Column Child4 (4, 3)
+ Column Parent2 (5, 2)
- + Column Child1 (5, 3)
- + Column Child2 (6, 3)
- + Column Child3 (7, 3)
- + Column Child4 (8, 3)
> res[1:4,]
         UNKNOWN
                                       UNKNOWN
                                                  UNKNOWN
                          UNKNOWN
                                                                    UNKNOWN
1 Column Parent1 Row Super-Parent Row Parent1 Row Child1 Row Child1-Child1
2 Column Parent1 Row Super-Parent Row Parent1 Row Child1 Row Child-Child2
3 Column Parent1 Row Super-Parent Row Parent1 Row Child2 Row Child-Child1
4 Column Parent1 Row Super-Parent Row Parent1 Row Child2 Row Child-Child2
  Column Child1 Column Child2 Column Child3 Column Child4
2
             13
                           23
                                          33
                                                         43
                                                         44
3
             14
                           24
                                          34
4
             15
                           25
                                          35
                                                         45
```

5.2 Reconstruction - Low Level Functions

The Low Level Reconstruction functions are called by the Main Reconstruction function. In particular, ReconsRowLabels is always called first and its results are one of the arguments for ReconsColLabels, which finishes the reconstruction of the entire LongForm Dataframe.

```
25a \langle Recons\ Low\ Level\ 25a \rangle \equiv 
\langle Recons\ Row\ Labels\ 25b \rangle 
\langle Recons\ Column\ Labels\ 26a \rangle
```

Uses rowvecs 23c.

5.2.1 Reconstruction - Row Labels

ReconsRowLabels iterates down the row parentage list (plist) recursively, extracting the names and using this to construct the columns of the finished LongForm Dataframe corresponding to the row labels. The final output is what was shown in the *Reconstruct Row Labels* chunk above as rowvecs[1:4,].

5.2.2 Reconstruction - Column Labels

As with the row labels, ReconsColLabels iterates down the column parentage list (plist) recursively. We also need to handle the parents differently from the lowest level child. The final output is what was shown in the *Reconstruct Col Labels* chunk above as res[1:4,].

```
⟨Recons Column Labels 26a⟩≡
ReconsColLabels =
function(plist, datamat, rowvecs){
   ⟨Recons Col Parents 26b⟩
   ⟨Recons Col Children 27⟩
   datfr
}
Defines:
   ReconsColLabels, used in chunks 24 and 26b.
Uses rowvecs 23c.
```

Example values for ToyExComplete.csv

```
> res[1:4,]
```

26a

```
UNKNOWN
                          UNKNOWN
                                       UNKNOWN
                                                  UNKNOWN
                                                                    UNKNOWN
1 Column Parent1 Row Super-Parent Row Parent1 Row Child1 Row Child1-Child1
2 Column Parent1 Row Super-Parent Row Parent1 Row Child1 Row Child-Child2
3 Column Parent1 Row Super-Parent Row Parent1 Row Child2 Row Child-Child1
4 Column Parent1 Row Super-Parent Row Parent1 Row Child2 Row Child-Child2
  Column Child1 Column Child2 Column Child3 Column Child4
             12
                            22
                                          32
2
             13
                            23
                                          33
                                                         43
3
             14
                            24
                                          34
                                                         44
4
             15
                            25
                                          35
                                                         45
```

Any parents are used to construct additional columns of factors (the labels of the parents) for the LongForm Dataframe, which is attached to the portion previously constructed in ReconsRowLabels.

Uses ReconsColLabels 26a and rowvecs 23c.

26

For the lowest level child, we extract the relevant 'data bits' from the original table and bind it to our Dataframe, using the lowest level child as the labels of these columns of data values.

Chunk Index 6

```
⟨Back End 10⟩
\langle Call\ Ident\ Main\ 23b \rangle
\langle Call\ Ident\ MostCommonBoundary\ 12c \rangle
⟨Call Pare Front 18b⟩
⟨Case Misaligned Col Parent 17c⟩
(Case Multiple Empty Rights 21)
(Case Single Empty Right 20e)
(Check for Other Empty Rights 20d)
(Collapse Fullrow Labels 18a)
(Construct rowslist and colslist 13a)
\langle Discern\ Parentage\ 16b \rangle
\langle document\ header\ 9a \rangle
\langle Front \ End \ 9b \rangle
\langle Generate\ Pattern\ vector\ 14a \rangle
(Get Non empty rows and cols 12b)
\langle Group \ by \ Pattern \ 14c \rangle
⟨Group Column Labels 13b⟩
(Ident by Most Common Boundary 12a)
(Ident Low Level 14d)
\langle Ident\ Main\ 11b \rangle
(Ident Most Common Boundary 16a)
(Ident Non Empty 15a)
\langle Ident\ Pattern\ 15b \rangle
\langle Identification 11a \rangle
\langle If first \ cell \ empty \ 19c \rangle
(If first column empty 19a)
\langle If \ only \ one \ column \ 18d \rangle
\langle If \ only \ one \ row \ 19b \rangle
(Otherwise call Pare Low Levels 20a)
(Pare By Empty Below 22a)
⟨Pare By Empty Right 20c⟩
\langle Pare\ Col\ 17b \rangle
\langle Pare\ Front\ 17a \rangle
\langle Pare\ Low\ Level\ 20b \rangle
\langle Pare\ Main\ 18c \rangle
\langle Recons\ Col\ Children\ 27 \rangle
⟨Recons Col Parents 26b⟩
⟨Recons Column Labels 26a⟩
(Recons Low Level 25a)
⟨Recons Main 23a⟩
\langle Recons \ Row \ Labels \ 25b \rangle
\langle Reconstruct\ Col\ Labels\ 24 \rangle
\langle Reconstruct Row Labels 23c \rangle
\langle Reconstruction 22b \rangle
\langle Table To Long Form. R \rangle
⟨ Take Largest Pattern 14b⟩
```

Identifier Index 7

Numbers indicate the chunks in which the function appears. Underline indicates the chunk where the function is defined.

attrLoc: 10, 18d, 19b, 20e, 21, 22a datacols: 17b, 17c, 18a, 18b, <u>23b</u>, 24

datarows: 23b, 23c, 24

IdentMain: 11b, 23b

IdentMostCommonBoundary: 12c, 16a

 $\begin{array}{lll} {\tt IdentNonEmpty:} & 12b,\,\underline{15a},\,23c\\ {\tt IdentPattern:} & 14a,\,\underline{15b}\\ {\tt labelcols:} & \underline{23b},\,23c \end{array}$

labelrows: 17b, 17c, 18a, 18b, 23b, 24

PareByEmptyBelow: $20a, \underline{22a}$ PareByEmptyRight: $20a, \underline{20c}$

PareCol: <u>17b</u>, 24

 $\begin{array}{lll} \texttt{PareFront:} & \underline{17a}, 18b, 23c \\ \texttt{PareMain:} & 17a, \underline{18c}, 19a, 20a \\ \texttt{ReconsColLabels:} & 24, \underline{26a}, 26b \\ \texttt{ReconsRowLabels:} & 23c, \underline{25b} \end{array}$

 $\verb"rowplist: \underline{23c},\,24$

rowvecs: 23c, 24, 25b, 26a, 26b, 27

 ${\tt TCRsink:} \ \ \underline{10}, \, 11b, \, 12c, \, 13a, \, 14a, \, 18d, \, 19b, \, 20e, \, 21, \, 22a, \, 23c, \, 24$

References

Davies, T., Bawa, Z., 2012a. The Journal of Community Informatics 8 (2). URL http://ci-journal.net/index.php/ciej/issue/view/41

Davies, T., Bawa, Z., 2012b. The Promises and Perils of Open Government Data (OGD). The Journal of Community Informatics 8 (2).

 ${\rm URL\ http://ci-journal.net/index.php/ciej/article/view/929}$

New Zealand Qualifications Authority, 2012. Secondary school statistics.

 $\label{lem:url:lem:u$

R Core Team, 2012. R: A Language and Environment for Statistical Computing. R Foundation for Statistical Computing, Vienna, Austria, ISBN 3-900051-07-0. URL http://www.R-project.org/

Ramsey, N., Sept 1994. Literate programming simplified. IEEE Software 11 (5), 97–105. URL http://www.cs.tufts.edu/~nr/noweb/

Statistics New Zealand, 2013. Infoshare.

URL http://www.stats.govt.nz/infoshare/